Signal Seeker

Tuesday, November 10, 2026

Boston · 43°F, mostly sunny · High 55° · Low 42° · Tomorrow: cloudy, 51°/40°

Nebius reports Thursday, Nov. 12.

IN THIS EDITION

Front Page 1

X 1
- Top of your feed 1
- From Your X 2
- The GOATs 3
- Investments 4

Evergreen 4

News 5

The Passage 5

Games 5

FRONT PAGE

Anthropic launches its IPO roadshow

Anthropic said it has launched the roadshow for its proposed initial public offering; its registration statement is filed with the SEC but not yet effective. @ParadisLabs flagged an $81B annual revenue run-rate as of Oct. 31, up $10B (14.1%) from September, while @jukan05 noted the company is pressing ahead even with its valuation cut in half from $2 trillion. @jimcramer called it the most consequential deal of our time, one the Jensen Huang story is obscuring. X @AnthropicAI · X @ParadisLabs · X @jukan05 · X @jimcramer

Report: Huang relationship tied to at least $7.5B in Nvidia chip shipments

Citing The Information, @KobeissiLetter reported that Jensen Huang's relationship with an alleged Chinese spy, said to date back to 2023, facilitated at least $7.5 billion in Nvidia chip shipments. A day earlier @GavinSBaker argued that since mid-2025 it has increasingly been Beijing, not Washington, blocking Nvidia GPUs in China. X @KobeissiLetter · X @GavinSBaker

SoftBank's CFO denies it defaulted on data-center turbine orders

Asked about Elon Musk's claim last week that SoftBank was defaulting on turbine orders, CFO Yoshimitsu Goto said it is "simply not correct at all. There is no such default," @aleabitoreddit reported, noting GE's CEO has also said he sees no cancellations. @rdd147 had argued the reported defaults spelled trouble for Anthropic's compute commitments. X @aleabitoreddit · X @rdd147

Musk: Starlink's network software turns out to suit training superintelligence

Elon Musk wrote that the software built to move vast amounts of data through space and down to the ground is also very well suited to training super intelligence, his most-engaged post in the feed this morning. X @elonmusk

X

TOP OF YOUR FEED

The most-engaged posts from your home timeline in the last 24 hours, a few per day and no more.

DK **Daniel Koss** @daniel_koss · Nov 10

I'm surprised to see that some $HIMS investors are not happy about this deal with Andrew Huberman.

You can argue about the price, but I think it's strategically incredibly smart and makes a ton of sense.

Here's why:

My $HIMS thesis is really simple: AI will lead to an explosion in healthcare solutions and innovation.

Where does the value accrue? I believe:

1\) Obviously in the companies that make these breakthroughs first and manage to protect their IP via patents.

BUT, that's obviously insanely hard to predict. You essentially have to gamble or have insider information. ...

> **mon** @moninvestor
> $HIMS - What a great addition to the team. @hubermanlab has 8 million subscribers on YouTube.

x.com/daniel_koss/status/2120058446850494608

TH **The Humanoid Hub** @TheHumanoidHub · Nov 10

Archon Robotics introduced Anchor 1, a whole-body controller for humanoid manipulation.

It turns sparse human targets (pelvis, wrists, ankles) into precise, stable robot motion, skipping online joint-space retargeting to avoid a common source of error.

It also separates motion encoding from feedback correction, pairing accurate tracking with smooth movement.

On a real robot, Anchor 1 posts the lowest tracking errors and highest task success against baselines including MimicLite, ScaleBFM, and SONIC, and collects teleoperation data more efficiently. …

x.com/TheHumanoidHub/status/2120042235819401394

J **Jukan** @jukan05 · Nov 10

This really isn't it...

I hope Mitsui speeds up its expansion.

This is far too late. China is going to devour market share. https://t.co/LyGjO1CoLV

x.com/jukan05/status/2120044654435676221

TF **The Fundamentalist** @Funmentalist · Nov 10

$NBIS

I saw the words “primarily” in the brackets and I immediately lost attention.

Learn to filter out noise, it will save you a lot of time! https://t.co/IZ5RKqJwa6

> **Brian Fry** @brianfry01
> $NBIS chose the asset light approach (primarily) and has done a fantastic job with that model so far but in my 28 years in HPC data center…

x.com/Funmentalist/status/2120081162156527645

MR **Milk Road AI** @MilkRoadAI · Nov 10

Kioxia is up about 8% today, but the stock is still down roughly 75% from its June peak.

For its third quarter, Kioxia expects NAND bit shipments to be flat and average selling prices to rise by a low to mid single digit percentage.

The fourth quarter is where it gets interesting. Kioxia sees bit shipments rising by high single digits to around 10% as enterprise SSD demand picks up.

Management also says NAND supply should stay tight, supported by long term customer agreements.

And Hong Kong just listed the world's first leveraged product on Kioxia today. …

> **Milk Road AI** @MilkRoadAI
> According to Korean media, South Korea's SSD exports hit $5.64 BILLION in October. That is up 584% from a year ago, and the ninth straight…

x.com/MilkRoadAI/status/2120066937996951784

M **mash** 🕶️ @Mhr1036 · Nov 10

It is increasingly clear that defensive cyber ability is the most important factor for choosing your inference, sandbox, training and other vendors.

Condolences for the amazing Fireworks team going through this tough time.

Maybe a defensive cyber coalition is on the cards https://t.co/Fo1gNXvLoR

x.com/Mhr1036/status/2120082569748197514

E **Every** 📧 @every · Nov 10

Fable 5.2 can work all weekend without you. Getting the most out of it means changing how you work.

.@kieranklaassen started building a video editor on a day off. Three days later, Everycut was live—with editing, visual effects, audio mixing, and an AI agent that could operate it. A week later, we ran a company-wide show-and-tell using it.

@hammer_mt gave it roughly 1,400 Slack messages and notes from 127 meetings. Without him outlining the argument first, it produced a cited account of his thinking, flagged 38 places where he'd contradicted himself, and made an explainer video. …

x.com/every/status/2120123331131355301

FROM YOUR X

AW **Alexandr Wang** @alexandr_wang · Nov 9

using muse ember🔥, each of these websites cost $0.04 on average

back in my day, these would cost you O($10,000) in custom dev shop costs

massive deflation in nice pretty websites https://t.co/l1EYh2kC27

x.com/alexandr_wang/status/2119866792348954839

FA **Fireside Alpha** @firesidealpha · Nov 9

JUST IN: Memphis Mayor reveals SpaceX data center taxes more than doubled to $54 million, after $25 million last year made it Shelby County's second largest taxpayer

Paul Young was asked about Elon Musk's claim that data centers have Memphis booming:

Q: “Recently, just speaking on the revenues the data centers are creating for Memphis. Elon Musk recently tweeted that

Memphis is 'booming' because of the data centers. He said that hotels and restaurants are full and called, tax revenue in Memphis tremendous right now. Do you just have any thoughts on that?"

YOUNG: "Well, we've talked a lot about the financial aspects of the project. I think that as we talk about data centers in our city, we want to talk about it in a balanced way. Of course there's an economic side. We want to make sure that we are articulating what protections and guardrails we're putting in place for the environmental sides that, people are concerned about, the quality of life, noise, traffic, all of those aspects."

"But, there there is tax revenue. I think this year it was $54,000,000 in taxes that have been generated. And I know last year when they generated about $25,000,000 in tax revenue, it was the second largest tax paying property in Shelby County. I don't know where $54,000,000 puts it, but it is significant taxes."

"And so as we look at, the tax revenue, we also wanna look at how do we ensure that there are proper protections. That's what this ninety day executive order has been about, and, it's the conversation that we'll continue to have with the community."

x.com/firesidealpha/status/2119934910723174807

THE GOATS

GB **Gavin Baker** @GavinSBaker · Nov 9

Something that merits more discussion amongst those who believe it was a mistake by both the Biden and Trump administrations to allow the export of less powerful GPUs to China:

Since the summer of 2025, it has increasingly been *China*, not just DC, blocking Nvidia GPUs in China. Within weeks of Secretary Lutnick saying the goal was to get Chinese developers "addicted" to the American technology stack, Beijing began pressuring its tech giants to stop buying the H20. Then it ordered them to stop buying the RTX Pro 6000D, sent customs teams to the ports to crack down on Nvidia chip imports, and barred foreign AI chips from new state-funded data centers. And as of last week, Reuters reports that roughly 600,000 H200s are sitting in Hong Kong because Beijing won't approve large-scale imports of a chip that is still better than anything China's own industry makes at scale.

If allowing the export of these less powerful GPUs was a policy mistake that *advantaged* China, then why is *China* the one blocking them?

From my perspective, China understands that if they are dependent on American AI hardware - either on the ground in China or accessed remotely via the wide range of southeast Asian clouds (which include SE Asian availability zones from the largest American hyperscalers) - then we have immense leverage over them in any future conflict. Being able to turn off the clouds instantly is actually the greatest leverage.

Yes, Chinese AI labs still want Nvidia chips. Gray market prices for H200 servers have roughly tripled over the past year. But that is exactly the point. Beijing is willing to impose real costs on its own AI labs to avoid dependence on American GPUs, which tells you how Beijing views GPU imports: not as a gift, but as a threat to its own ambitions.

Maybe, just maybe, selling China GPUs that trail our own state of the art, so that Chinese developers stay "addicted" to American GPUs and China doesn't build an alternative AI hardware ecosystem that we have no control or influence over *and* that China could share over time with our adversaries, was actually a smart, pro-American national defense policy. And letting them access southeast Asian clouds might be an even smarter policy as so easy to turn clouds off.

I believed and still believed it was smart policy. And I say this as someone whose posts here are hawkish enough on China that I won't travel there anymore.

And note that both the Biden and Trump administration allowed the export of less powerful GPUs to China for all of these reasons. It was only under the Trump administration that China woke up to the risk of what was actually a pretty smart, pro-national defense policy and then China began restricting GPUs.

The lesson of the H20 is that a deliberately crippled chip doesn't keep China on our stack. Beijing moved against the H20 and RTX Pro 6000D as it concluded that domestic Chinese chips had caught up to them. So if China develops an alternative AI hardware ecosystem that it shares with Russia, Iran and North Korea, we will look back and think it would have been even smarter to offer China the B30A in 2025 instead of the H200s we eventually approved case-by-case. Especially given the H200 is not the frontier: with Vera Rubin now in full production, it is two generations behind what American labs are buying.

All of this is to say that Jensen has been consistently lobbying for a position that was adapted by both the Biden *and* Trump admins *and* looks smarter every month that goes by as Chinese labs increasingly adapt their MoE architectures to domestic Chinese hardware. DeepSeek reportedly gave Huawei and Cambricon early access to V4 instead of Nvidia and AMD, and a Huawei-led team has since post-trained V4-Pro on Ascend chips.

And if America *ever* falls behind China in AI, the China hawks here in America who helped accelerate the creation of this alternative Chinese AI hardware ecosystem will deserve a significant share of the blame.

x.com/GavinSBaker/status/2119908820608798745

Chris Camillo @ChrisCamillo · Nov 9

Imagine building a $4 trillion empire to get taken down by a situationship

x.com/ChrisCamillo/status/2119850141142389078

INVESTMENTS

Serenity @aleabitoreddit · Nov 10

If you look at CW laser ASPs for AI data centers, the pricing differences are pretty substantial:

70mW: ~$3–4
100mW: ~$4–6 (Source Photonics channel checks).

With some premiums apparently going up to $7-$9 range with the recent shortage for pluggables.

120–150mW NPO: ~$15–20
350mW UHP CW: ~$40 (Nomura)
400mW UHP CW for NPO/CPO: ~$50–75 (Mizuho)

With former communication laser chips going from 20-50 cents.

The pricing and gross margins of these players go up with each architectural shift (60%+ for CPO lasers).

And go up even more depending how far down the stack you go (selling ELSFP was ~2.5x revenue content)

In terms of how I see the landscape now:

- $COHR is probably the best risk adjusted laser player now, for people who don't like uncertainty. Their new $20B LTA is very substantial.

- $LITE pretty negative now given the valuation overhang from $MRVL acquisition. Operationally very opsitive.

- $AAOI has extreme demand visibility (likely $6-$9B in new LTAs from just 1 quarter), but credible execution uncertainty. They've also been focusing on pluggables + bottlenecked by DSPs, rather focusing on UHP CPO laser merchant markets.

But I think if markets see a new $5B-$7B LTA that they're in late stage discussion for... and see the $4.4B MC. Pretty sure that's a catalyst for recovery.

The demand doesn't disappear but timeline shifts back a few months.

AOI also signaled DSP as the main bottleneck, last quarter it was TIAs, so other different component players that are interesting.

- $SIVE has extreme capacity expansion (300m laser capacity targeted in total).

The economic implications that they won't disclose are likely enormous relative to their MC, even if they sell $3 70mW lasers to China pluggable players.

But there's a qualification uncertainty to fill the gap, that markets have punished the company for. Especially when management communication amplifies uncertainty, despite the ongoing shortage.

Right now entire AI trade is going up and down together though.

x.com/aleabitoreddit/status/2120074819878449605

M. V. Cunha @mvcinvesting · Nov 9

🚨 $RKLB just announced a record quarter for Neutron launch contracts.

Eight missions across five operators, totaling over $515M, with several launches priced at $70M each.

One of those customers is planning to deploy data centers in space. https://t.co/rbJFS21VS0

x.com/mvcinvesting/status/2119901996941750278

EVERGREEN

POST · NOV 7

I'm going to give a contrarian take on Neoclouds like $NBIS

Serenity @aleabitoreddit

After the valuation crash, Serenity argues neoclouds with Nvidia allocations, contracted and permitted capacity, hyperscaler customers and sane financing will do fine, pointing to SpaceX's $40B GPU financing and IREN's 90% prepayments and recontracted H100s at double prior pricing. The real risk, in this view, sits with newcomers still raising money or waiting on permits.

x.com/aleabitoreddit/status/2119077473107472729

POST · OCT 20

Been obsessed with the stock market since I was a teenager

Chris Camillo @ChrisCamillo

A veteran of the dotcom bubble and the financial crisis says he has never seen skepticism like the one around AI: each bubble narrative is replaced by another, and easily debunked claims get repeated by media and analysts. He expects spectacular failures along the way, but finds it hard to understand why so many investors seem to be rooting for AI to fail.

x.com/ChrisCamillo/status/2112632034497630228

NEWS

AI data centres are hunting for used turbines and jet engines. With new gas turbines backordered to 2030, AI data-centre builders are buying used turbines and converted jet engines. Reuters (via Investing.com)

Disney and Adobe will build custom AI models for the parks, studio and ESPN. Disney and Adobe signed a multi-year deal to build private generative AI models trained on Disney's IP. Reuters (via Investing.com)

FDA adds a rare-blindness warning to all GLP-1 drugs. The FDA now requires a warning on all GLP-1 drugs about NAION, a rare cause of sudden vision loss. NBC News

An mRNA treatment got heart muscle to regrow after heart attacks in mice. A Cambridge mRNA therapy made mouse heart cells divide after a heart attack, improving pumping by about 10%. University of Cambridge

THE PASSAGE

FROM YOUR OLD QUOTE BANK

You cannot overestimate the unimportance of practically everything.

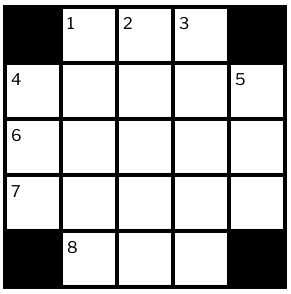
— John Maxwell

GAMES

MENTAL MATH

1. 17 × 9 = ?
2. 35% of 960 = ?
3. Fermi: a sheet of printer paper is about 0.1 mm thick. How tall is a stack of 40 reams of 500 sheets?

MINI CROSSWORD



Across

1 Make a sad face and thrust out one's lower lip (3)
4 Someone who drives racing cars at high speeds (5)
6 A particular environment or walk of life (5)
7 Any creature of exceptional size (5)
8 One or some or every or all without specification (3)

Down

1 A dark region of considerable extent on the surface of the moon (5)
2 A large body of water constituting a principal part of the hydrosphere (5)
3 A coin worth one-hundredth of the value of the basic unit (5)
4 A small piece of cloth or paper (3)
5 Any of various long-tailed rodents similar to but larger than a mouse (3)

FOUR CARDS

If a card has a vowel on one side, it has an even number on the other.

E	K	3	8

Each card has one thing on each side. Which cards must you turn over to check the rule?

RIDDLE

What has hands but cannot clap?

End of today's paper.

Covering roughly Nov 9–10, 2026. Compiled Tuesday morning, ET. Links are in the web edition.

ANSWERS

MENTAL MATH

1. **153** 17 × 9 = 10 × 9 + 7 × 9 = 153
2. **336** 10% is 96, so 35% = 3 × 96 + half of 96 = 336
3. **about 2 m** 40 × 500 = 20,000 sheets × 0.1 mm = 2,000 mm = 2 m

MINI CROSSWORD

	1 M	2 O	3 P	
4 R	A	C	E	5 R
6 A	R	E	N	A
7 G	I	A	N	T
	8 A	N	Y	

FOUR CARDS

Turn E and 3. E could hide something that breaks the rule; 3 breaks it if the other side turns out to be E. 8 can't break it (the rule never says only those cards have it), and K isn't covered by the rule at all. Most people pick the card that matches the rule's second half instead of the one that could falsify it.

RIDDLE

A clock.